

Lecture 2: Sampling-based Approximations And Function Fitting

Yan (Rocky) Duan
Berkeley AI Research Lab

Many slides made with John Schulman, Xi (Peter) Chen and Pieter Abbeel

Quick One-Slide Recap

■ Optimal Control

=

gamma between 0-1
H finite or infinite

given an MDP (S, A, P, R, γ, H)

find the optimal policy π^*

■ Exact Methods:



Value Iteration



Policy Iteration

Limitations:

- Update equations require access to dynamics model
- Iteration over / Storage for all states and actions: requires small, discrete state-action space

-> **sampling-based approximations**

-> **Q/V function fitting**

Sampling-Based Approximation

- Q Value Iteration
- Value Iteration?
- Policy Iteration
 - Policy Evaluation
 - Policy Improvement?

Recap Q-Values

$Q^*(s, a)$ = expected utility starting in s , taking action a , and (thereafter) acting optimally

Bellman Equation:

$$Q^*(s, a) = \sum_{s'} P(s'|s, a)(R(s, a, s') + \gamma \max_{a'} Q^*(s', a'))$$

Q-Value Iteration:

$$Q_{k+1}(s, a) \leftarrow \sum_{s'} P(s'|s, a)(R(s, a, s') + \gamma \max_{a'} Q_k(s', a'))$$

(Tabular) Q-Learning

- Q-value iteration: $Q_{k+1}(s, a) \leftarrow \sum_{s'} P(s'|s, a) (R(s, a, s') + \gamma \max_{a'} Q_k(s', a'))$
- Rewrite as expectation: $Q_{k+1} \leftarrow \mathbb{E}_{s' \sim P(s'|s, a)} \left[R(s, a, s') + \gamma \max_{a'} Q_k(s', a') \right]$
- (Tabular) Q-Learning: replace expectation by samples
 - For an state-action pair (s, a) , receive: $s' \sim P(s'|s, a)$
 - Consider your old estimate: $Q_k(s, a)$
 - Consider your new sample estimate:
$$\text{target}(s') = R(s, a, s') + \gamma \max_{a'} Q_k(s', a')$$
 - Incorporate the new estimate into a running average:
$$Q_{k+1}(s, a) \leftarrow (1 - \alpha) Q_k(s, a) + \alpha [\text{target}(s')]$$

(Tabular) Q-Learning

Algorithm:

Start with $Q_0(s, a)$ for all s, a .

Get initial state s

For $k = 1, 2, \dots$ till convergence

 Sample action a , get next state s'

 If s' is terminal:

$$\text{target} = R(s, a, s')$$

 Sample new initial state s'

 else:

$$\text{target} = R(s, a, s') + \gamma \max_{a'} Q_k(s', a')$$

$$Q_{k+1}(s, a) \leftarrow (1 - \alpha)Q_k(s, a) + \alpha [\text{target}]$$

$$s \leftarrow s'$$

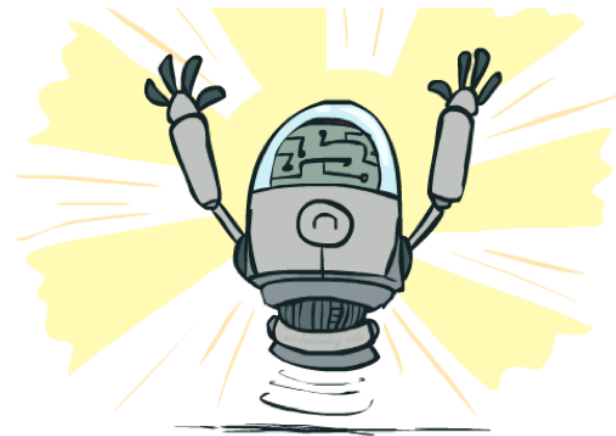
two sampling processes for next state s' , and new initial state

How to sample actions?

- Choose action that maximizes $Q_k(s, a)$ (i.e. greedily)?
- Choose random actions?
- ϵ -Greedy: choose random action with prob. ϵ , otherwise choose action greedily

Q-Learning Properties

- Amazing result: Q-learning converges to optimal policy -- even if you're acting suboptimally!
- This is called **off-policy learning**
- Caveats:
 - You have to explore enough
 - You have to eventually make the learning rate small enough
 - ... but not decrease it too quickly

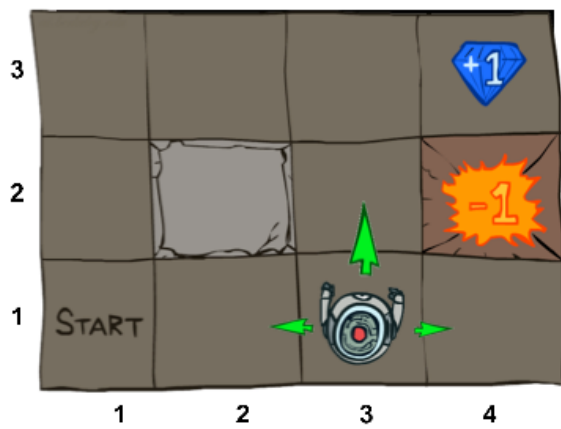


Q-Learning Properties

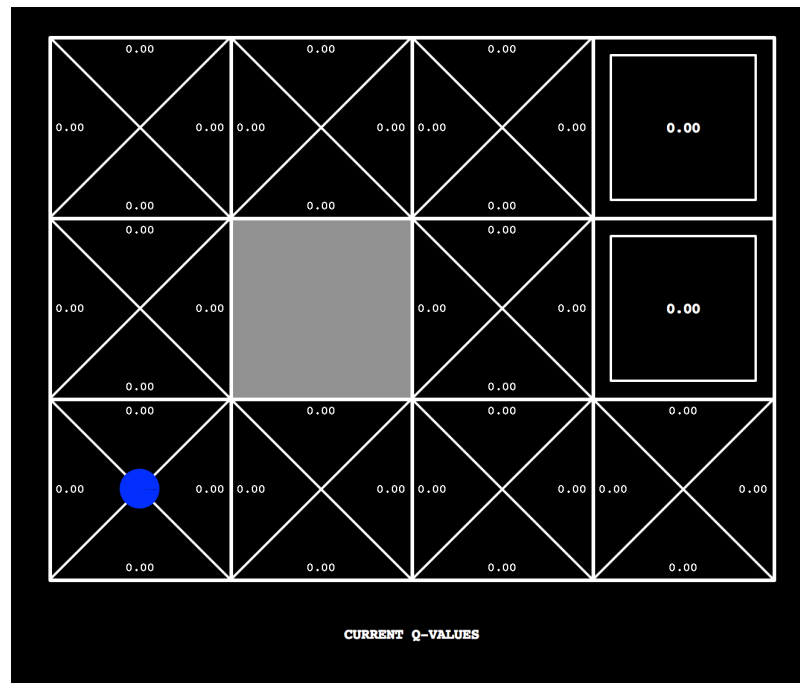
- Technical requirements.
 - All states and actions are visited infinitely often
 - Basically, in the limit, it doesn't matter how you select actions (!)
 - Nonnegative learning rates, and for all state and action pairs (s,a):

$$\sum_{t=0}^{\infty} \alpha_t(s, a) = \infty \qquad \sum_{t=0}^{\infty} \alpha_t^2(s, a) < \infty$$

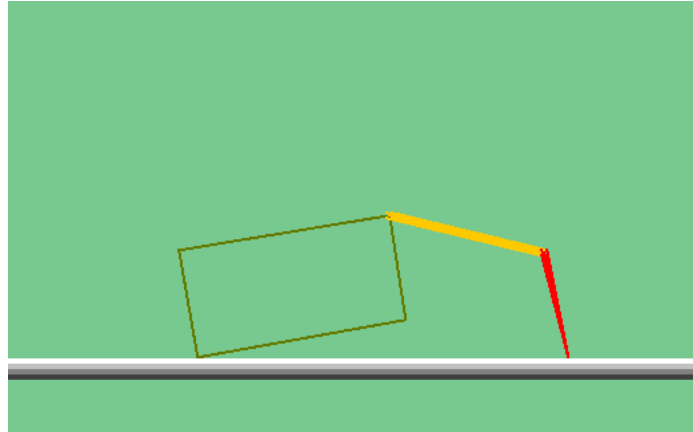
Q-Learning Demo: Gridworld



- States: 11 cells
- Actions: {up, down, left, right}
- Deterministic transition function
- Learning rate: 0.5
- Discount: 1
- Reward: +1 for getting diamond, -1 for falling into trap



Q-Learning Demo: Crawler



- **States:** discretized value of 2d state: (arm angle, hand angle)
- **Actions:** Cartesian product of {arm up, arm down} and {hand up, hand down}
- **Reward:** speed in the forward direction

Sampling-Based Approximation

✓ Q Value Iteration → (Tabular) Q-learning

- Value Iteration?
- Policy Iteration
 - Policy Evaluation
 - Policy Improvement?

Value Iteration w/ Samples?

- Value Iteration

$$V_{i+1}^*(s) \leftarrow \max_a \mathbb{E}_{s' \sim P(s'|s,a)} [R(s, a, s') + \gamma V_i^*(s')]$$

- unclear how to draw samples through max.....

Sampling-Based Approximation

✓ Q Value Iteration → (Tabular) Q-learning

■ ~~Value Iteration?~~

■ Policy Iteration

■ Policy Evaluation

■ Policy Improvement?

Recap: Policy Iteration

One iteration of policy iteration:

- Policy evaluation for current policy π_k :

- Iterate until convergence

$$V_{i+1}^{\pi_k}(s) \leftarrow \sum_{s'} P(s'|s, \pi_k(s)) [R(s, \pi_k(s), s') + \gamma V_i^{\pi_k}(s')]$$

Can be approximated by samples

This is called Temporal Difference (TD) Learning

- Policy improvement: find the best action according to one-step look-ahead

$$\pi_{k+1}(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V^{\pi_k}(s')]$$

Unclear what to do with the max (for now)

Sampling-Based Approximation

- ✓ ■ Q Value Iteration → (Tabular) Q-learning
- ~~Value Iteration?~~
- Policy Iteration
 - ✓ ■ Policy Evaluation
 - ~~Policy Improvement (for now)~~

Quick One-Slide Recap

- Optimal Control

=

given an MDP (S, A, P, R, γ, H)

find the optimal policy π^*

- Exact Methods:



Value Iteration



Policy Iteration

Limitations:

- Update equations require access to dynamics model
- Iteration over / Storage for all states and actions: requires small, discrete state-action space

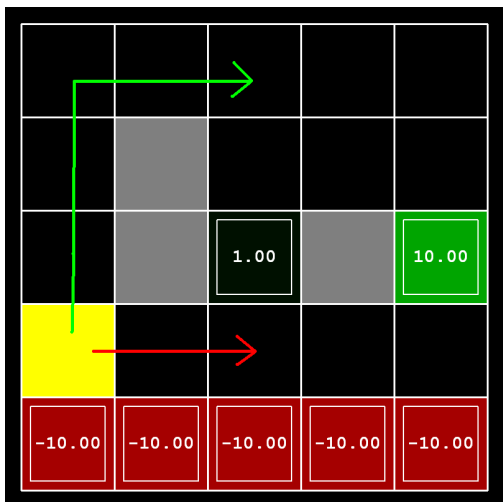


sampling-based approximations

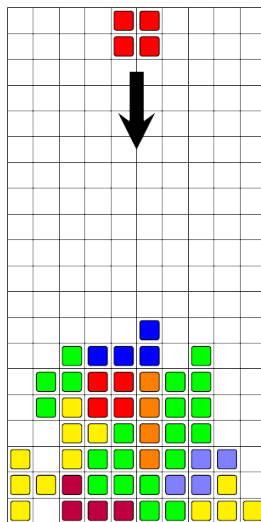
-> Q/V function fitting

Can tabular methods scale?

- Discrete environments



Gridworld
 10^1



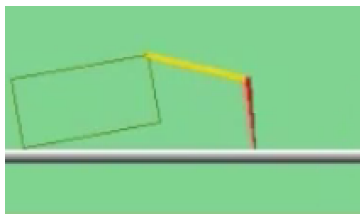
Tetris
 10^{60}



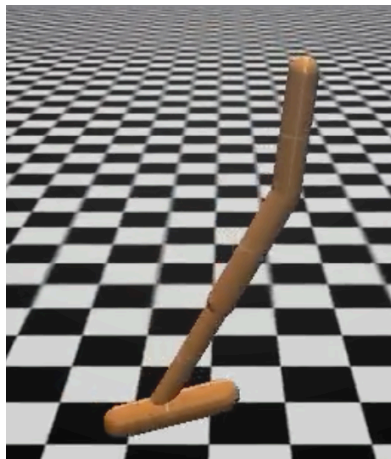
Atari
 10^{308} (ram) 10^{16992} (pixels)

Can tabular methods scale?

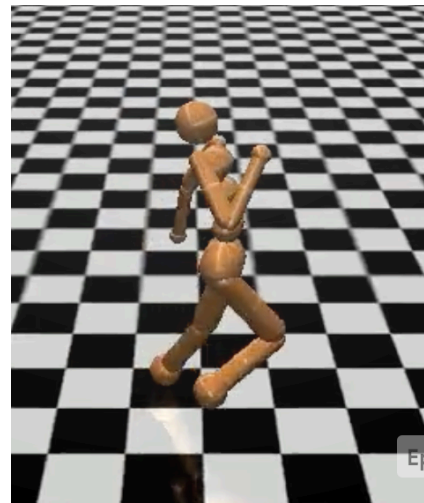
- Continuous environments (by crude discretization)



Crawler
 10^2



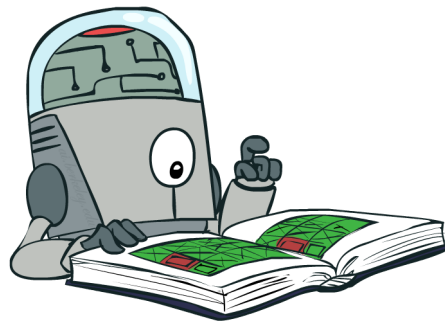
Hopper
 10^{10}



Humanoid
 10^{100}

Generalizing Across States

- Basic Q-Learning keeps a table of all q-values
- In realistic situations, we cannot possibly learn about every single state!
 - Too many states to visit them all in training
 - Too many states to hold the q-tables in memory
- Instead, we want to generalize:
 - Learn about some small number of training states from experience
 - Generalize that experience to new, similar situations
 - This is a fundamental idea in machine learning, and we'll see it over and over again



Approximate Q-Learning

- Instead of a table, we have a parametrized Q function: $Q_\theta(s, a)$

- Can be a linear function in features:

$$Q_\theta(s, a) = \theta_0 f_0(s, a) + \theta_1 f_1(s, a) + \cdots + \theta_n f_n(s, a)$$

- Or a complicated neural net

- Learning rule:

- Remember: $\text{target}(s') = R(s, a, s') + \gamma \max_{a'} Q_{\theta_k}(s', a')$

- Update:

$$\theta_{k+1} \leftarrow \theta_k - \alpha \nabla_\theta \left[\frac{1}{2} (Q_\theta(s, a) - \text{target}(s'))^2 \right] \Big|_{\theta=\theta_k}$$

Connection to Tabular Q-Learning

- Suppose $\theta \in \mathbb{R}^{|S| \times |A|}$, $Q_\theta(s, a) \equiv \theta_{sa}$

$$\begin{aligned} & \nabla_{\theta_{sa}} \left[\frac{1}{2} (Q_\theta(s, a) - \text{target}(s'))^2 \right] \\ &= \nabla_{\theta_{sa}} \left[\frac{1}{2} (\theta_{sa} - \text{target}(s'))^2 \right] \\ &= \theta_{sa} - \text{target}(s') \end{aligned}$$

- Plug into update: $\theta_{sa} \leftarrow \theta_{sa} - \alpha(\theta_{sa} - \text{target}(s'))$
 $= (1 - \alpha)\theta_{sa} + \alpha[\text{target}(s')]$


- Compare with Tabular Q-Learning update:

$$Q_{k+1}(s, a) \leftarrow (1 - \alpha)Q_k(s, a) + \alpha [\text{target}(s')]$$

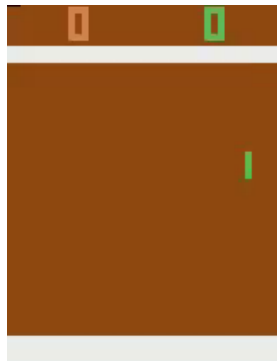
Engineered Approximation Example: Tetris

- state: naïve board configuration + shape of the falling piece $\sim 10^{60s}$ states!
- action: rotation and translation applied to the falling piece
- 22 features aka basis functions ϕ_i
 - Ten basis functions, $0, \dots, 9$, mapping the state to the height $h[k]$ of each column.
 - Nine basis functions, $10, \dots, 18$, each mapping the state to the absolute difference between heights of successive columns: $|h[k+1] - h[k]|$, $k = 1, \dots, 9$.
 - One basis function, 19, that maps state to the maximum column height: $\max_k h[k]$
 - One basis function, 20, that maps state to the number of 'holes' in the board.
 - One basis function, 21, that is equal to 1 in every state.

$$\hat{V}_\theta(s) = \sum_{i=0}^{21} \theta_i \phi_i(s) = \theta^\top \phi(s)$$



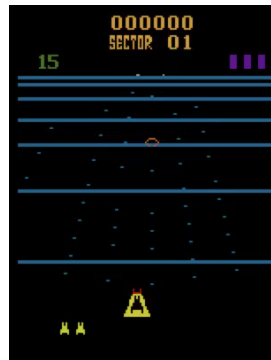
Deep Reinforcement Learning



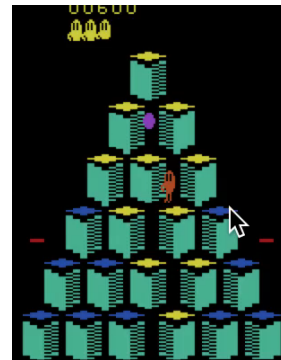
Pong



Enduro



Beamrider



Q*bert

- From pixels to actions
- Same algorithm (with effective tricks)
- CNN function approximator, w/ 3M free parameters

Lab 1

- We have now covered enough materials for Lab 1.
- Will be released on Piazza by this afternoon.
- Covers value iteration, policy iteration, and tabular Q-learning.